

Per-Step Inference Latency vs. Real-Time Budget

